


Journal Pre-proof

Extracting declarative constraints for process modeling from natural language descriptions with large language models

Gyunam Park, Julian Kofferath, Minsu Cho



PII: S0167-9236(26)00016-3

DOI: <https://doi.org/10.1016/j.dss.2026.114627>

Reference: DECSUP 114627

To appear in: *Decision Support Systems*

Please cite this article as: G. Park, J. Kofferath and M. Cho, Extracting declarative constraints for process modeling from natural language descriptions with large language models, *Decision Support Systems* (2026), doi: <https://doi.org/10.1016/j.dss.2026.114627>.

This is a PDF of an article that has undergone enhancements after acceptance, such as the addition of a cover page and metadata, and formatting for readability. This version will undergo additional copyediting, typesetting and review before it is published in its final form. As such, this version is no longer the Accepted Manuscript, but it is not yet the definitive Version of Record; we are providing this early version to give early visibility of the article. Please note that Elsevier's sharing policy for the Published Journal Article applies to this version, see: <https://www.elsevier.com/about/policies-and-standards/sharing#4-published-journal-article>. Please also note that, during the production process, errors may be discovered which could affect the content, and all legal disclaimers that apply to the journal pertain.

© 2026 Published by Elsevier B.V.

Extracting Declarative Constraints for Process Modeling From Natural Language Descriptions with Large Language Models

anonymous authors

Abstract

With the growing availability of unstructured text data in organizations, automating the extraction of process models from natural language descriptions has become increasingly crucial. Traditional rule-based techniques face challenges such as limited generalization and constraints to narrow sets of control-flow constructs. We present a systematic empirical study of Large Language Models (LLMs) for translating natural language sentences into declarative process constraints using the *Declare* modeling language. Through comprehensive evaluation across seven LLM architectures, we demonstrate that fine-tuned models achieve 97.8% template accuracy compared to 53.8% for existing rule-based approaches. Our fine-tuning approach substantially outperforms prompting techniques (99.4% vs. 56.5% template accuracy), establishing clear guidance for practical deployment. We contribute a benchmark dataset of 969 sentence-constraint pairs across 11 *Declare* templates, direction-sensitive evaluation metrics, and complete reproducibility materials. The extracted constraints enable practical decision support applications including compliance monitoring, conformance checking, and automated governance dashboards.

Keywords: Process Modeling, Automated Extraction, Large Language Models

¹ 1. Introduction

² Process models are foundational to effective process management, serving as
³ critical tools for understanding, analyzing, and optimizing business operations [1].
⁴ These models provide a structured representation of business activities, facilitating
⁵ more precise communication and better compliance with standards. Accurately
⁶ mapping out processes can significantly enhance operational efficiency and agility.

However, process actors and domain experts lack the expertise required to construct process models. Moreover, the process of eliciting these models can be exceedingly time-consuming [2]. Additionally, process-related information is often documented in more accessible formats, such as text documents [2, 3].

11 Recognizing this, researchers have turned to Natural Language Processing (NLP)
12 approaches to extract process knowledge, especially process models, from textual
13 sources [2, 3, 4, 5, 6, 7, 8, 9, 10]. This approach offers a promising way to convert
14 unstructured textual data into formal process representations.

However, as described in [11, 12], the research field is fragmented and lacks systematic efforts. First, existing rule-based methods often lack flexibility, leading to poor generalization across varied text inputs. Second, existing methods support only a narrow range of control-flow constructs, limiting their applicability in complex process scenarios. Third, the field lacks a unified framework with robust datasets and evaluation metrics, which hinders consistent assessment and comparison of methods.




Figure 1: An illustration of the proposed method and a possible use case.

In this paper, we present a method to discover Declare process models from natural language descriptions using LLMs. First, as shown in Figure 1, a fine-tuned

23 LLM translates each sentence in the natural language description into a Declare
 24 constraint such as *Precedence(Patient Registration, Medical Examination)* and *Re-*
 25 *sponse(Examination, Prescription Issuance)*. Next, to build a cohesive Declare
 26 process model, we match activity names, e.g., *Examination* to *Medical Examination*.
 27 Finally, we merge Declare constraints to construct a Declare process model. The
 28 resulting process model can be used in various ways; it can serve as a benchmark
 29 for conducting declarative conformance checking against a de facto process model
 30 derived from event data [13].

31 Our approach contributes to the automated process extraction as follows:

32 1. *Reliable component for template translation:* By utilizing LLMs, our approach
 33 significantly improves generalization capabilities and accuracy over rule-based
 34 methods for sentence-to-constraint translation. The component adapts to
 35 diverse sentence structures and linguistic variations and supports 11 *Declare*
 36 templates.

37 2. *Benchmark dataset with provenance and policy:* We provide a benchmark-
 38 quality dataset of 969 textual descriptions paired with ground-truth *Declare*
 39 constraints across 11 templates, together with detailed provenance, annotation
 40 guidelines (including ambiguity policy), and intended use for benchmarking.

41 3. *Evaluation design and baselines:* We evaluate with separate metrics for tem-
 42 plate classification and activity extraction (direction-sensitive for ordered
 43 templates), and compare against both prompting and a rule-based state of the
 44 art, demonstrating higher reliability of the fine-tuned component.

45 4. *Reproducibility materials:* We release scripts and configuration to reproduce
 46 training, evaluation, and inference, along with hardware setup and runtime
 47 characteristics, enabling transparent replication.

48 This work focuses on a core component within a larger pipeline that transforms
 49 natural language documents into *Declare* constraints and models. The component
 50 translates a single sentence into a single constraint and currently assumes a one-
 51 sentence-to-one-constraint mapping. We discuss implications of this assumption
 52 and outline integration with upstream and downstream steps, such as sentence
 53 segmentation, multi-constraint detection and splitting (e.g., for conjunctions), and

54 activity matching, to form coherent models.

55 The remainder of this paper is organized as follows: In Section 2, we introduce
 56 related work concerning the synthesis of process models from natural language.
 57 Section 3 introduces Declare and decoder-only LLMs. Section 4 and Section 5 de-
 58 scribe the extraction pipeline and employed techniques. Next, Section 6 evaluates the
 59 performance of our approach and compares it with existing methods. In Section 7, we
 60 present a detailed discussion of our findings, including limitations and opportunities
 61 for future research. Finally, Section 8 concludes this paper.

62 **2. Related Work**

63 Recently, translating natural language texts to process models has gained attention.
 64 Bellan et al. [11] compare ten state-of-the-art approaches to extracting process models
 65 from text, highlighting the need for common benchmarks and robust evaluation.

66 Classical NLP-based approaches translate text to imperative models (e.g., BPMN):
 67 De A. R. Gonçalves et al. [5] (storytelling mining), Friedrich et al. [6] (semantic
 68 parsing), Text Process Miner [7], Honkisz et al. [8], and Sholiq et al. [9]. Beyond
 69 full model discovery, works extract specific elements (activities, actors, gateways)
 70 [14, 15, 16, 17] and align text with existing models [18, 19, 20].

71 For declarative models, López et al. [4] use DCR graphs via BERT-based clas-
 72 sification and NER, while Van der Aa et al. [2] propose an NLP pipeline to extract
 73 *Declare* constraints (five templates). Prompting LLMs has also been explored: Grohs
 74 et al. [21] show GPT-4 can outperform the pre-LLM baseline on several templates.

75 More recently, LLM-based methods have been applied to constraint extraction
 76 from documents and semantics-aware process mining. Barrientos et al. [22] and
 77 Mustroph et al. [23] translate compliance requirements from text to check them over
 78 event logs. Rebmann et al. [24] assess LLM capabilities for semantics-aware tasks
 79 in process mining, while Busch et al. [25] introduce explainable semantic anomaly
 80 detection using sequence-to-sequence models.

81 In parallel, conversational and imperative modeling with LLMs is emerging:
 82 Köpke and Safan [26] present efficient conversational process modeling; Klievtsova

et al. [27] survey conversational process modeling applications and implications.

Connections between imperative and declarative representations are relevant to our setting. Weidlich and Van der Werf [28] study relational semantics for Petri nets; Barbaro et al. [29] demonstrate deriving LTLf declarative specifications from sound workflow nets. Such mappings suggest that text-to-imperative datasets could be repurposed to derive text-to-declarative pairs in future work.

Finally, foundational LLM reports (e.g., Gemma, LLaMA, Mistral) and instruction tuning/PEFT literature underpin our design choices for fine-tuning and inference [30].

Our work focuses on a reusable component for sentence-to-*Declare* constraint translation across 11 templates, intended to plug into full NL-to-constraints pipelines. It complements (i) document-level constraint extraction and semantics-aware mining (broader granularity), (ii) conversational and imperative modeling (different artifacts), and (iii) imperative-to-declarative mappings (useful to enrich datasets). In this sense, our contribution lies at the intersection of fine-grained sentence translation and broader process modeling efforts.

Despite progress, prior work exhibits recurring limitations: (1) rule-based methods struggle to generalize across diverse linguistic structures; (2) many approaches support only a narrow subset of control-flow constructs; and (3) the field lacks standardized evaluation frameworks and benchmark datasets for consistent comparison. Our study addresses these gaps by (a) fine-tuning LLMs to improve generalization and direction-sensitive template recognition, (b) supporting a broad set of 11 *Declare* templates, and (c) contributing a benchmark dataset with clear annotation policy and an evaluation protocol that separates template and activity metrics.

3. Preliminaries

In this work, we design an approach to the automated extraction of *Declare* process models using *decoder-only LLMs*. In this section, we first introduce the process modeling language *Declare* and then explain the architecture and functionality of decoder-only LLMs.

112 3.1. Declare

113 Declare is a process modeling language built upon Linear Temporal Logic (LTL)
 114 that provides templates for expressing constraints on single or multiple activities. The
 115 templates originate from common temporal logic patterns used in model checking [31]
 116 and are categorized into occurrence patterns (e.g., Absence, Existence) and order
 117 patterns (e.g., Precedence, Response) [32].

Table 1: A selected set of common Declare templates.

Constraint	Textual	Graphical
Init(A)	A is the first to occur	
End(A)	A is the last to occur	
AtLeastOne(A)	A occurs at least once	
AtMostOne(A)	A occurs at most once	
Response(A,B)	If A occurs, then B occurs after A	
Precedence(A,B)	If B occurs, then A occurs before B	
RespondedExistence(A,B)	If A occurs, then B occurs before or after A	
ChainResponse(A,B)	If A occurs, then B occurs immediately afterwards	
ChainPrecedence(A,B)	If B occurs, then A occurs immediately beforehand	
CoExistence(A,B)	If A occurs, then B occurs before or after A, and vice versa	
Succession(A,B)	A occurs if and only if it is followed by B	
NotCoExistence(A,B)	A and B never occur together	
NotSuccession(A,B)	B cannot occur after A	


118 Given its focus on business processes, Declare interprets LTL specifications over
 119 finite-length process executions. Table 1 shows common Declare templates with their
 120 natural language descriptions and graphical notation.

121 **Definition 1 (Declare Constraint).** Let \mathcal{A} be the universe of activity names. A
 122 Declare constraint is characterized by the pair $(k(x_1, \dots, x_m), \kappa)$, where:

- 123 • $k(x_1, \dots, x_m)$ is a Declare template;
- 124 • κ is a mapping that associates each variable x_i (for $i \in \{1, \dots, m\}$) with an
 125 activity $\kappa(x_i) = a_i \in \mathcal{A}$.

126 Such a Declare constraint can be represented as $k(a_1, \dots, a_m)$ for brevity.

127 A Declare process model is a collection of Declare constraints defined over
 128 consistent activity names. Figure 2 depicts an example of a declare model.


**Figure 2:** A graphical Declare model.

129 **Definition 2 (Declare Process Model).** A Declare process model is a tuple $DS =$
130 (Act, K) , comprising:

- 131 • $Act \subseteq \mathcal{A}$ is a set of activities;
132 • K is a set of Declare constraints defined over Act ;

133 **3.2. Decoder-only LLMs**

134 We adopt decoder-only Transformer architectures [33] trained with a causal
135 language modeling objective: given a textual prompt, the model autoregressively
136 predicts the next token under a causal mask. A high-level view is shown in Figure 3.
137 For prompting baselines, we leverage instruction-following techniques (e.g., Chain-of-
138 Thought and zero-shot reasoning) [34, 35], and note that self-consistency can further
139 stabilize reasoning [36].

**Figure 3:** An overview of the LLM architecture. **(a)** shows the overall decoder-only model structure. **(b)** zooms into one decoder block, consisting of self-attention (SA), and feed-forward network (FFN) layers.

¹⁴⁰ **4. Fine-Tuning Large Language Models**

¹⁴¹ This section outlines fine-tuning LLMs to translate text inputs to Declare con-
¹⁴² straints. Figure 4 demonstrates how the fine-tuned LLM is incorporated into the over-
¹⁴³ all extraction pipeline. In the following, we first describe the downstream task for fine-
¹⁴⁴ tuning LLMs and then explain the fine-tuning process based on the downstream task.




Figure 4: The overall pipeline required to transform a process description into a Declare model.

¹⁴⁵ *4.1. Downstream Task*

¹⁴⁶ The first step in the extraction pipeline involves passing each sentence from the
¹⁴⁷ input process description separately to the fine-tuned LLM to obtain the formal
¹⁴⁸ Declare constraint described by this sentence.

¹⁴⁹ The fine-tuning process employs *causal language modeling*, where the model is
¹⁵⁰ trained to predict the next token in a sequence given the previous tokens. This method
¹⁵¹ is particularly suitable for generation tasks, e.g., translating sentences into constraints,
¹⁵² as it allows the model to learn the sequential nature of the output. The model gener-
¹⁵³ ates constraints token-by-token, using the previously generated tokens and the input
¹⁵⁴ sentence context. During inference, the number of generated tokens can be bounded.

155 Additionally, a custom formatting function preprocesses the training data for the
 156 model. This function creates input sequences where the sentence (textual descrip-
 157 tion) and its expected constraint are concatenated, ensuring that the model receives
 158 well-prepared inputs aligned with the task requirements. This leads to improved
 159 performance and generalization capabilities. The format is as follows:

160 Question: What is the Declare constraint described in the following
 161 sentence? <description>

162 Answer: <gold standard>

163 During inference, we prompt the fine-tuned model using the same format, prepend-
 164 ing the question before the current sentence. For each inference iteration, i.e., for each
 165 input sentence, we obtain a response from the model that adheres to the above format.

166 To extract formal Declare constraints from the LLM’s response at inference time,
 167 we use a regular expression pattern to search within the decoded output. This pattern
 168 is designed to match all considered types of Declare constraints (templates), such
 169 as *Precedence*, *Init*, *AtLeastOne*, etc., followed by activities enclosed in parentheses.
 170 By leveraging the fixed text-based format of Declare constraints, we can accurately
 171 identify and extract these constraints from the model’s responses. Additionally, by
 172 defining the output format using the formatting function and limiting the number
 173 of generated tokens, we ensure that there is, at most, one match for the pattern in
 174 the decoded output.

175 4.2. Fine-Tuning

176 To adopt a decoder-only large language model (LLM) like Google’s Gemma-7b for
 177 the task of translating input sentences into formal Declare constraints, we employ Low-
 178 Rank Adaptation (LoRA) [37]. LoRA is a parameter-efficient fine-tuning technique
 179 that modifies the LLM architecture by introducing low-rank adaptation matrices
 180 into the model’s transformer layers, allowing efficient task-specific fine-tuning.

181 With LoRA, the fine-tuning process involves adding small, trainable low-rank
 182 matrices to the self-attention and feed-forward components of the model. These
 183 matrices adjust the query, key, and value projections in the self-attention mechanism,
 184 as well as the intermediate layers in the feed-forward network. Specifically, instead

185 of directly updating the large weight matrices, LoRA decomposes these updates into
 186 products of smaller matrices, significantly reducing the number of parameters that
 187 need to be fine-tuned. Figure 5 illustrates the affected components.




Figure 5: (a) zooms into one decoder block of the LLM, consisting of self-attention (SA), and feed-forward network (FFN) layers. (b) shows the detailed structure of the LoRA adaptation. The components shown in grey are kept frozen during fine-tuning, whereas the orange components contain updated weights.

188 For a given pre-trained weight matrix $W_0 \in \mathbb{R}^{d \times k}$, LoRA introduces two trainable
 189 weight matrices, $W_{up} \in \mathbb{R}^{d \times r}$ and $W_{down} \in \mathbb{R}^{r \times k}$, where the rank r is much smaller
 190 than the minimum of d and k . Normally, given input h_{in} , the output is computed
 191 as $h_{out} = W_0 \cdot h_{in}$. LoRA modifies this output by introducing an incremental update
 192 ΔW that encapsulates task-specific knowledge:

$$h_{out} = W_0 \cdot h_{in} + \alpha \cdot \Delta W \cdot h_{in} = W_0 \cdot h_{in} + \alpha \cdot W_{up} \cdot W_{down} \cdot h_{in},$$

193 where α is a scaling factor [30].

194 *Training Configuration and Implementation Details.* To fine-tune the large language
 195 model, we carefully select hyperparameters to optimize performance and accuracy.
 196 The training is conducted with a batch size of 2 per device and gradient accumulation
 197 steps set to 4, ensuring sufficient exposure to data and effective gradient updates.
 198 We include a warmup phase of 2 steps to stabilize the initial training and limit the

total training steps to 200 to prevent overfitting. The learning rate is set at 2×10^{-4} to ensure a balanced update of the model weights, promoting convergence without overshooting the optimal parameter values.

For reproducibility, we provide complete implementation details:

- *Models and tokenizers*: We consider `google/gemma-7b` [38], `mistralai/Mistral-7B-v0.3` [39], `tiuae/falcon-7b` [40], `meta-llama/Llama-2-7b-hf` [41], `meta-llama/Llama-3-8b` [42], and small efficient baselines `meta-llama/Llama-3.2-1B-Instruct` and `meta-llama/Llama-3.2-3B-Instruct` [42]. Unless stated otherwise, we use base (non-instruct) variants and the corresponding official tokenizers (snapshots from October 2024 for 7B/8B models; August 2025 for 1B/3B).
- *Fine-tuning setup*: Parameter-efficient fine-tuning via PEFT/LoRA [30, 37] (`peft==0.8.2`) with rank $r = 8$ and scaling $\alpha = 8$, dropout at default; target modules include `q_proj`, `k_proj`, `v_proj`, `o_proj`, `gate_proj`, `up_proj`, `down_proj` for attention and MLP projections. Optimizer: AdamW [43] (Transformers default betas/eps; weight decay default); learning rate 2×10^{-4} with linear decay and warmup = 2 steps; per-device batch size 2 with gradient accumulation 4 (effective 8); total training steps 200 (single epoch given dataset size).
- *Inference and parsing*: Greedy decoding (temperature 0.0, `max_new_tokens=15`). Outputs are post-processed with a regex that matches one *Declare* template name followed by activities in parentheses; only the first valid match is retained.
- *Hardware*: NVIDIA H100 96GB VRAM (single GPU) on a compute cluster; CPU/RAM standard server-class (not critical for reproduction).
- *Artifacts*: Code and dataset: https://anonymous.4open.science/r/dss_rev/.

Hardware and runtime. All fine-tuning experiments were executed on a single NVIDIA H100 (96 GB VRAM). With the LoRA configuration above, per-fold fine-tuning of a 7B model for 200 steps required approximately 10–20 minutes; complete 10-fold cross-validation finished within 2–4 hours per model. Typical inference

²²⁹ throughput was \sim 2–5 ms per sentence on GPU (batch size 1) and \sim 10–30 ms per
²³⁰ sentence on a modern CPU, excluding I/O.

²³¹ **5. Activity Matching**

²³² This section describes the second step in extracting a coherent Declare process
²³³ model (cf. Figure 4). We first introduce the problem that this step aims to tackle
²³⁴ and then present several solutions to the problem.

²³⁵ *5.1. Problem Definition*

²³⁶ So far, in the process of extracting a process model from text, we have isolated
²³⁷ Declare constraints, each extracted separately from a single sentence. To create
²³⁸ a coherent process model, it is crucial that activities with the same meaning are
²³⁹ labeled consistently across all constraints. In other words, we need to ensure that
²⁴⁰ semantically equivalent activities share identical labels across all constraints.

²⁴¹ To formalize the requirement of ensuring consistent labeling, we describe a ded-
²⁴² icated function. Such a function takes as input a set of Declare constraints K and
²⁴³ outputs Declare process model $DS = (Act, K')$ with a modified set of constraints K'
²⁴⁴ where semantically equivalent activities have identical labels, and Act is the set of
²⁴⁵ activity equivalence classes in K .

²⁴⁶ Formally, K' is obtained by modifying K such that:

$$K' = \{(k(\vec{x}), \kappa') | (k(\vec{x}), \kappa) \in K, \kappa' = \text{UpdateMapping}(\kappa, Act)\}$$

²⁴⁷ The function $\text{UpdateMapping}(\kappa, Act)$ performs activity matching by replacing
²⁴⁸ each activity label in constraint mapping κ with its canonical representative from the
²⁴⁹ corresponding equivalence class in Act . Given a constraint with activities mapped
²⁵⁰ by κ , the function identifies which equivalence class each activity belongs to and
²⁵¹ substitutes it with the chosen canonical label for that class, ensuring consistent
²⁵² labeling across all constraints while preserving the original constraint semantics.

253 Consider the following set of constraints with slightly different surface forms for
 254 the same activities:

$$K = \{Response(UserLogin, EnterMenu), Precedence(UploadFile, DeleteFile), \\ Response(FileDel, ProcessFile), ChainResponse(ProcessingFile, ApplyChange)\}$$

255 The activity matching process identifies semantically equivalent activities and groups
 256 them into equivalence classes:

$$Act = \{[UserLogin], [EnterMenu], [UploadFile], \\ [DeleteFile \equiv FileDel], [ProcessFile \equiv ProcessingFile], [ApplyChange]\}$$

257 where each class contains semantically equivalent activity labels, with the canonical
 258 representative listed first. The UpdateMapping function replaces each activity label
 259 with its canonical representative, yielding:

$$K' = \{Response(UserLogin, EnterMenu), Precedence(UploadFile, DeleteFile), \\ Response(DeleteFile, ProcessFile), ChainResponse(ProcessFile, ApplyChange)\}$$

260 *5.2. Solutions*

261 We provide two solutions to the presented problem.

262 *Zero-Shot Prompting.* First, a zero-shot prompting approach prompts LLMs by
 263 providing only instructions for the desired output without any examples of how
 264 the task has been performed correctly [35]. Full prompt templates are provided in
 265 Appendix A. Note that variations exist in how activities are expressed, ranging
 266 from verb-based (e.g., *DeleteFile*) to noun-based (e.g., *FileDel*) and gerund-based
 267 (e.g., *ProcessingFile*) formulations.

268 *Few-Shot Prompting.* Another approach is to provide examples of input with correct
 269 output. Such a prompt, where the input includes a few examples to guide the model
 270 in performing a specific task, is called a *few-shot* prompt [35]. Full few-shot prompt
 271 is in Appendix B.

272 We include the full list of constraints in the prompt because matching decisions
 273 often rely on cross-constraint evidence. Providing complete constraint context re-
 274 duces ambiguity and promotes consistent canonical labels compared to processing
 275 constraints in isolation. Activities alone lose critical relational and directional signals
 276 needed for correct canonicalization. For example, from activities-only $\{SubmitForm,$
 277 $ReviewRequest, ApproveRequest, ArchiveCase\}$, a matcher may erroneously merge
 278 $ReviewRequest$ and $ApproveRequest$ into a single $Request$ label. However, the full
 279 constraints $\{Response(SubmitForm, ReviewRequest), Precedence(ApproveRequest,$
 280 $ArchiveCase)\}$ reveal distinct roles: $ReviewRequest$ follows submission while $Ap-$
 281 $proveRequest$ must precede archiving, indicating they should remain separate. Sim-
 282ilarly, constraint context enables proper cross-constraint consolidation, identifying
 283 that $FileDel$ and $DeleteFile$ represent the same concept, as do $ProcessingFile$ and
 284 $ProcessFile$, allowing unification to canonical forms.

285 6. Evaluation

286 This section aims to validate the efficiency of our proposed method both inter-
 287 nationally and externally. The experimental results reported in the evaluation are fully
 288 reproducible at https://anonymous.4open.science/r/dss_rev/.

289 6.1. Benchmark Dataset

290 To support the development and evaluation of our approach, we prepared a
 291 benchmark dataset comprising 969 text descriptions paired with corresponding
 292 gold-standard Declare constraints across a variety of domains, including healthcare,
 293 corporate compliance, education, and finance¹. The corpus mixes curated domain sen-
 294 tences and synthetic paraphrases to ensure lexical and structural diversity. Sampling
 295 targeted balanced coverage over the 11 templates, with additional paraphrase variants
 296 for common templates (e.g., $Response$, $Precedence$). Figure 6 illustrates the distribu-
 297 tion of 11 Declare tempaltes, with $Response$ and $Precedence$ being the most prevalent
 298 due to their central role in capturing process dependencies and order relations.

¹The dataset is publicly available at https://anonymous.4open.science/r/dss_rev/data




Figure 6: Distribution of Declare constraints in the benchmark dataset

299 *Annotation guidelines and decision rules.* Annotators were instructed to (i) extract
 300 verb- or noun-based activity mentions as concise Pascal-style labels (e.g., *Upload-*
 301 *File*, *FileDel*); (ii) map each text to the best-fitting *Declare* template considering
 302 order constraints (e.g., if B must follow A, prefer *Response(A,B)* over *RespondedEx-*
 303 *istence(A,B)* when ordering is explicit); and (iii) avoid conflating disjunctions or
 304 conjunctions into single activities unless explicitly intended. Canonicalization rules
 305 (e.g., verb-first preference) were applied consistently to reduce label drift.

306 *Ambiguity and multiple-valid-templates policy.* For under-specified sentences that ad-
 307 mit multiple valid templates, we adopted a balanced annotation strategy. When a sen-
 308 tence could reasonably be interpreted with multiple valid templates, annotators were
 309 instructed to select the least represented template in the dataset to promote class bal-
 310 ance. For example, "If a file is uploaded, it must be processed." could
 311 permit both *RespondedExistence(UploadFile, ProcessFile)* and *Response(UploadFile,*
 312 *ProcessFile)* depending on whether explicit ordering is assumed. In such cases, the
 313 less frequent template was chosen to mitigate dataset imbalance.

314 *Adjudication and agreement.* Each instance was labeled by one annotator and
 315 screened by a second.

316 6.2. Evaluation Metrics

317 In assessing the accuracy of the constraints generated by the fine-tuned LLM,
 318 it is essential to identify a suitable metric. For each sentence in the respective test

319 dataset, we have both the gold standard constraint and the discovered constraint.

320 There are three types of errors:

- 321 • *Template error*: the discovered template is incorrect
- 322 • *Activity error*: at least one of the extracted activities is incorrect
- 323 • *Direction error*: the extracted activities are correct, but the order is incorrect

324 The existing work evaluates the approach by framing the problem as a slot-filling
 325 problem [2]. They assume that each constraint can be represented as a set of slots.
 326 For example, a binary constraint like $\text{Precedence}(X, Y)$ has three slots: one for the
 327 template and one for each activity involved. Although this method provides a straight-
 328 forward way to compute precision and recall, it conflates different types of errors.

329 When the gold standard constraint is $\text{Precedence}(X, Y)$ and the predicted con-
 330 straint is $\text{Precedence}(Y, X)$, the slot-filling approach assigns partial credit based on over-
 331 lapping slots. However, this error is not a simple activity mismatch; thus, it is a direc-
 332 tion error that fundamentally reverses the meaning of the constraint. Similarly, when
 333 comparing $\text{Precedence}(X, Y)$ with $\text{Response}(X, Y)$, the activities are identical but the
 334 templates differ, yet the slot-filling evaluation still grants a relatively high score, failing
 335 to capture the semantic gap between the two constraints. Therefore, it lacks the ability
 336 to separately capture template accuracy, activity correctness, and direction sensitivity.

337 To address these limitations and provide more precise evaluation, we introduce
 338 direction-sensitive metrics that properly distinguish template and activity accuracy.

339 Let the gold constraint be $g = k(a_1, \dots, a_m)$ and the prediction be $p = \tilde{k}(b_1, \dots, b_n)$.
 340 We evaluate template, activities, and direction with simple, direction-sensitive scores
 341 without weights:

$$342 \quad \text{TemplateAcc}_i = \mathbb{1}[\tilde{k} = k], \\ 343 \quad \text{ActAcc}_i = \begin{cases} \frac{1}{m} \sum_{j=1}^m \mathbb{1}[a_j = b_j], & \text{if } \tilde{k} = k \text{ and } n = m, \\ 0, & \text{otherwise,} \end{cases} \\ 344 \quad \text{ConstraintScore}_i = \text{TemplateAcc}_i \cdot \text{ActAcc}_i.$$

344 Thus, a wrong template yields zero score; a correct template yields partial credit
 345 for correctly placed activities (direction-sensitive by construction). For symmetric

³⁴⁶ templates (e.g., *CoExistence*, *NotCoExistence*), we evaluate ActAcc_i as the maximum
³⁴⁷ over argument permutations that preserve the template’s semantics.

³⁴⁸ Over a dataset of N instances, we report macro-averaged metrics. Template
³⁴⁹ accuracy is the fraction of instances with correct template prediction. Activity ac-
³⁵⁰ curacy is computed only for instances with correct templates, measuring the fraction
³⁵¹ of correctly positioned activities. Constraint score is the overall fraction of instances
³⁵² with both correct template and correct activities.

³⁵³ 6.3. Internal Validation

³⁵⁴ We internally validate which LLM architecture performs best for sentence-to-
³⁵⁵ constraint translation through systematic comparison across seven open-source,
³⁵⁶ decoder-only LLMs. We run 10-fold cross-validation on our benchmark dataset to
³⁵⁷ identify the optimal model for external validation and analyze performance patterns
³⁵⁸ across different model families and sizes.

³⁵⁹ 6.3.1. Experimental Setup

³⁶⁰ A 10-fold cross-validation was utilized to evaluate our approach. The benchmark
³⁶¹ dataset was randomly partitioned into ten equally sized, non-overlapping subsets.
³⁶² Each subset was used once as a test set, while the LLM was fine-tuned using the other
³⁶³ nine subsets as training data. For each model, the training parameters were identical,
³⁶⁴ as described in Subsection 4.2, ensuring fair comparison across architectures.

³⁶⁵ 6.3.2. Experimental Results

³⁶⁶ Table 2 reports macro-averaged TemplateAcc, conditional ActAcc, and the overall
³⁶⁷ ConstraintScore as defined in Subsection 6.2.

³⁶⁸ Gemma 7B delivers the strongest performance across all three metrics, achieving
³⁶⁹ 97.8% template accuracy, 78.2% conditional activity accuracy, and 76.4% overall
³⁷⁰ constraint score. This establishes it as our primary model for subsequent external
³⁷¹ validations, with a substantial 9.5 percentage point advantage in template accuracy
³⁷² over the next-best performing Llama-3 8B. Clear generational improvements emerge
³⁷³ within model families, most prominently between Llama variants. Llama-3 8B
³⁷⁴ significantly outperforms Llama-2 7B (88.3% vs. 66.8% template accuracy).

Table 2: Cross-validation performance of different fine-tuned decoder-only LLMs.

Model	TemplateAcc	ActAcc (cond.)	ConstraintScore
Gemma 7B (ft)	0.978	0.782	0.764
Llama-3 8B (ft)	0.883	0.752	0.664
Falcon 7B (ft)	0.880	0.755	0.664
Mistral 7B (ft)	0.740	0.780	0.577
Llama-2 7B (ft)	0.668	0.733	0.491
Llama-3.2 3B-Instruct (ft)	0.920	0.701	0.644
Llama-3.2 1B-Instruct (ft)	0.885	0.664	0.588

375 An interesting pattern emerges when comparing template classification versus
 376 activity extraction performance across all models. Template accuracy consistently
 377 falls below conditional activity accuracy, indicating that LLMs are generally proficient
 378 at identifying relevant activities within text but struggle more with distinguishing
 379 between semantically similar constraint types, such as Response versus Respond-
 380 edExistence. Mistral 7B exhibits a particularly unusual pattern with relatively low
 381 template accuracy (74.0%) but high conditional activity accuracy (78.0%).

382 The instruction-tuned lightweight models demonstrate competitive performance,
 383 particularly in template classification tasks. Llama-3.2 1B-Instruct achieves 88.5%
 384 template accuracy, which exceeds several 7B-8B base models. The 3B variant per-
 385 forms even better with 92.0% template accuracy, nearly matching the performance
 386 of much larger models while requiring significantly fewer computational resources.

387 6.4. External Validation - Part A

388 In this section, we conduct external validation by comparing our approach to
 389 the state-of-the-art method by Van der Aa et al. [2]. We use the best-performing
 390 model from the internal validation (Gemma 7B fine-tuned with LoRA as specified in
 391 Subsection 4.2). The comparison focuses on the four templates supported by both
 392 approaches: *Init*, *End*, *Precedence*, and *Response*. We evaluate performance on two
 393 distinct validation sets to assess generalization and reliability improvements over a
 394 traditional rule-based NLP pipeline.

395 *6.4.1. Experimental Setup*

396 We limit the evaluation to *Init*, *End*, *Precedence*, and *Response*, which are the
 397 templates supported by the baseline. We exclude *Succession*, as it is a conjunction
 398 of *Response* and *Precedence*.

399 For fine-tuning the best-performing model (Gemma 7B), we constructed a train-
 400 ing set of 850 sentence-constraint pairs by starting from our benchmark dataset and
 401 removing all sentences that occur in the external validation data of Van der Aa et
 402 al. [2] (including minor paraphrases), thereby avoiding any overlap or leakage. No
 403 sentence from *V1* or *V2* is included in training.

404 We report only template accuracy (TemplateAcc) in this comparison. Activity-
 405 level metrics are not directly comparable due to differing activity vocabularies and
 406 canonicalization policies between approaches.

407 To assess performance, we use a human-checked validation set (*V1*) consisting
 408 of 80 inputs covering the four templates, disjoint from the training data. In addition,
 409 we validate both approaches on a subset of the validation set used by Van der Aa
 410 et al. [2], denoted *V2*. From the original 103 sentences, we exclude those with
 411 *Succession* or negative templates as gold standards, and those that do not describe
 412 one of the four templates (e.g., pure conditional constructs).

413 *6.4.2. Experimental Results*

414 As shown in Table 3, the best-performing fine-tuned model (Gemma 7B) sub-
 415 stantially improves TemplateAcc on both *V1* and *V2*. We report TemplateAcc only
 416 by design, since activity strings and canonicalization differ across systems and would
 417 confound activity-level comparisons. In our validation sets, gold labels use concise,
 418 verb-based Pascal-case activities; under this metric choice, the fine-tuned model
 419 attains nearly perfect template accuracy on *V1* (versus 53.8% for the baseline) and
 420 improves over the baseline on *V2* as well.

421 *6.5. External Validation - Part B*

422 Next, we compare our approach with prompting baselines. Prior work [21] shows
 423 that strong prompting can outperform traditional rule-based systems [2]. Here we

⁴²⁴ evaluate whether a small, fine-tuned model (Gemma 7B with LoRA) provides more
⁴²⁵ reliable performance across our 11 templates.

⁴²⁶ *6.5.1. Baselines*

⁴²⁷ *Few-shot prompting.* One-shot per template (11 shots in total), each mirroring the
⁴²⁸ abstract template description in Table 1. The prompt asks to output the constraint
⁴²⁹ only (no prose). A sketch of the prompt is given in Appendix C.

⁴³⁰ *Few-shot Chain-of-Thought (CoT) prompting.* One-shot per template with short
⁴³¹ rationales preceding the answer, following Wei et al. [34], Kojima et al. [35]. We use
⁴³² the instruction-tuned Gemma 7B-it for CoT and standard Gemma 7B for non-CoT.
⁴³³ An excerpt is shown in Appendix D.

⁴³⁴ *Prompt formatting, token budgeting, and decoding.*

- ⁴³⁵ • *Formatting.* Prompts are rendered in monospace blocks; CoT uses conversation
⁴³⁶ control tokens required by Gemma 7B-it. When using other models, adapt
⁴³⁷ or remove control tokens to match the target chat template. We instruct the
⁴³⁸ model to output only a single constraint.

Table 3: Comparison with the state-of-the-art Van der Aa et al. [2]. Entries in the two rightmost columns specify TemplateAcc.

		Cases	State-of-the-art	Our approach
			[2]	(Gemma 7B)
V1	Init	11	0.273	0.917
	End	4	0.0	1.0
	Response	38	0.684	1.0
	Precedence	27	0.519	1.0
	Overall	80	0.538	0.988
V2	Init	4	0.75	1.0
	End	1	0.0	1.0
	Response	14	0.643	0.929
	Precedence	27	0.963	0.962
	Overall	45	0.822	0.956

- 439 • *Token budgeting.* To stay within context limits we (i) keep shots concise, (ii)
 440 prefer schema-like definitions over verbose prose, (iii) reuse short, canonical
 441 activity names, and (iv) trim boilerplate/whitespace. If close to the window,
 442 reduce the number of shots first, then shorten rationales. We verify that
 443 system/instruction + shots + input fit within the model’s maximum tokens.
- 444 • *Decoding parameters.* Decoding is greedy (temperature 0.0; no sampling;
 445 top-p=1.0; top-k disabled) with `max_new_tokens=15`. For CoT, we add stop
 446 sequences for end-of-turn tokens. Outputs are regex-validated to extract exactly
 447 one *Declare* constraint (template and activities) in canonical form.
- 448 • *Common failure cases.* Over-length prompts (truncated rationales), malformed
 449 outputs (multiple constraints/extraneous prose), and ambiguity-induced activity
 450 order swaps. Mitigations include stricter formatting, shorter shots, explicit
 451 stop sequences, and regex validation.
- 452 • *Common failure cases.* Over-length prompts (truncated rationales), malformed
 453 outputs (multiple constraints/extraneous prose), and ambiguity-induced activity
 454 order swaps. Mitigations include stricter formatting, shorter shots, explicit
 455 stop sequences, and regex validation.

456 6.5.2. Data Collection

457 We utilize the same dataset as for the internal validation. From this dataset, we
 458 allocate 80% as training data, while the remaining 20% serves as a test set for both
 459 the prompting approaches and the fine-tuned LLM.

460 6.5.3. Experimental Setup

461 We use Gemma-7b for both fine-tuning and the *few-shot* prompt. For the *few-*
 462 *shot* *CoT* approach, we utilize the instruct version of Gemma-7b. The prompts are
 463 designed to consider the same set of *Declare* templates on which we trained the
 464 LLM: *Init*, *End*, *AtMostOne*, *AtLeastOne*, *Response*, *Precedence*, *ChainResponse*,
 465 *ChainPrecedence*, *RespondedExistence*, *CoExistence*, and *NotCoExistence*.

466 6.5.4. Experimental Results

467 As shown in Table 4, fine-tuning yields the highest TemplateAcc (0.994) and the
 468 best overall *ConstraintScore* (0.809). CoT improves TemplateAcc over plain few-shot
 469 (0.565 vs. 0.459) and markedly boosts conditional *ActAcc* (0.734), indicating better
 470 activity placement when the template is correct. However, the fine-tuned model
 471 balances both template recognition and activity placement, leading to the strongest
 472 end-to-end constraint quality.

473 7. Discussion

474 7.1. Findings

475 *Model Architecture and Performance Patterns.* Our internal validation reveals clear
 476 performance hierarchies among LLM architectures for constraint extraction. Gemma
 477 7B consistently delivers superior performance across all metrics (97.8% template
 478 accuracy, 78.2% conditional activity accuracy, 76.4% constraint score), establish-
 479 ing a 9.5 percentage point advantage over the next-best Llama-3 8B. Surprisingly,
 480 model size does not universally correlate with performance. The compact Llama-3.2
 481 3B-Instruct (92.0% template accuracy) significantly outperforms the larger Llama-2
 482 7B, highlighting the effectiveness of instruction tuning and knowledge distillation.
 483 The 1B-Instruct variant achieves 88.5% template accuracy, exceeding several 7B-8B
 484 base models and suggesting that instruction tuning provides substantial benefits for
 485 structured classification tasks.

486 *Template vs. Activity Extraction Performance.* An interesting pattern emerges across
 487 all models: template accuracy consistently exceeds conditional activity accuracy,

Table 4: Prompting baselines versus fine-tuned model using the direction-sensitive metrics.

Technique	TemplateAcc	ActAcc (cond.)	ConstraintScore
Few-shot prompt	0.459	0.051	0.024
Few-shot CoT prompt	0.565	0.734	0.415
Our approach	0.994	0.814	0.809

488 indicating that LLMs excel at distinguishing constraint types but face greater chal-
 489 lenges in precise activity extraction and positioning. This suggests that template
 490 classification and activity extraction present distinct computational challenges, with
 491 template recognition being the more tractable task.

492 The few instances where fine-tuned LLMs fail to identify correct templates
 493 stem primarily from linguistic ambiguities inherent in natural language. For exam-
 494 ple, "Upon receiving a support ticket, the troubleshooting process begins"
 495 permits multiple valid interpretations: either *Init(ReceiveSupportTicket)* focusing on
 496 the initiation aspect, or *Response(ReceiveSupportTicket, BeginTroubleshooting)* em-
 497 phasizing the causal relationship. The desired activity abstraction level significantly
 498 influences template selection in such cases.

499 *Comparison with State-of-the-Art.* External validation demonstrates substantial im-
 500 provements over existing rule-based approaches. On validation set *V1*, our Gemma
 501 7B model achieves 98.8% template accuracy compared to 53.8% for Van der Aa et
 502 al. [2]. The largest performance gaps occur in unary constraints (*Init*, *End*), where
 503 rule-based methods struggle due to their reliance on detecting specific meta-actions
 504 and keywords like "process", "start", and "end". Our LLM-based approach
 505 overcomes these limitations through robust natural language understanding.

506 For common binary templates (*Response*, *Precedence*), the fine-tuned model
 507 demonstrates significant improvements, effectively learning to distinguish between
 508 semantically similar but structurally different constraint types. On validation set *V2*,
 509 our approach maintains competitive performance across most templates, with the ex-
 510 ception of *Precedence* constraints where the baseline achieves artificially high accuracy
 511 (96.3%) due to default constraint generation rather than sophisticated NLP processing.

512 *Fine-tuning vs. Prompting Strategies.* Comparison with prompting baselines confirms
 513 that parameter adaptation through fine-tuning is essential for reliable constraint ex-
 514 traction. While chain-of-thought prompting improves over basic few-shot approaches
 515 (56.5% vs. 45.9% template accuracy), both fall substantially short of fine-tuned
 516 performance (99.4% template accuracy). This indicates that prompting techniques,

despite their sophistication, cannot substitute for task-specific weight adaptation when precise structural understanding is required.

Advanced prompting provides interpretable reasoning steps and better activity placement when templates are correctly identified (73.4% conditional activity accuracy for CoT), but struggles with consistent template classification. The fine-tuned model achieves superior balance across both template recognition and activity extraction, leading to the highest end-to-end constraint quality.

7.2. Practical Implications for Decision Support Systems

Discovered *Declare* constraints provide machine-checkable policy abstractions that can be operationalized in decision support:

- *Compliance and governance*: Constraints act as executable rules derived from textual policies and can be monitored continuously on event data to detect violations and support governance dashboards.
- *Policy automation*: Violations or risk conditions can trigger proactive alerts and workflow actions, supporting policy enforcement and case-handling automation.
- *Audit readiness and explainability*: Findings can be linked to concrete trace evidence (case IDs, timestamps, activity pairs), yielding auditable, explainable outcomes grounded in the original text-derived rules.

7.3. Academic Implications

Our contributions also have implications for research methodology and future studies in IS, BPM, and NLP:

- *Benchmark and comparability*: The curated benchmark (969 sentence–constraint pairs across 11 templates) and simplified, direction-sensitive metrics (TemplateAcc, conditional ActAcc, ConstraintScore) enable transparent, apples-to-apples comparisons and reduce evaluation ambiguity.
- *Reproducibility and extensibility*: Released scripts, prompts, and configurations, together with hardware/runtime notes, facilitate faithful replication and lower the barrier for follow-up work, including re-running with new base models or decoding settings.

- 546 • *Cross-paradigm bridges*: The dataset and protocol can support studies mapping
- 547 between imperative and declarative representations, or repurposing text-to-
- 548 imperative resources to synthesize text-to-declarative pairs.
- 549 • *Model-size insights*: Including 1B/3B instruction-tuned baselines highlights
- 550 viable low-resource settings and opens avenues on efficiency–accuracy trade-offs,
- 551 PEFT choices, and deployment constraints.

552 *7.4. Limitations and Threats to Validity*

553 *Scope and expressiveness*. Our component assumes a one-sentence-to-one-constraint
 554 mapping; multi-constraint sentences and cross-sentence constraints require additional
 555 logic (e.g., splitting/aggregation). We currently do not cover advanced *Declare*
 556 constructs such as branched responses or time conditions, which limits expressiveness
 557 in complex scenarios. Generalization to highly domain-specific phrasing remains
 558 challenging where data are scarce.

559 *LLM-related threats*. As highlighted by recent work on contamination and evalua-
 560 tion malpractices [44], LLM-based studies risk inadvertent data leakage (train/test
 561 overlap) and protocol biases. We mitigate this by (i) careful dataset curation and
 562 cross-validation splits, (ii) disjoint external validations, (iii) transparent reporting
 563 of models, prompts, and configs, and (iv) releasing code and data for reproducibility.
 564 Prompt overfitting is reduced by fine-tuning and by keeping decoding simple (greedy)
 565 with regex post-processing.

566 *Model-size scope*. Our primary comparisons focus on 7B–8B models; we addi-
 567 tionally include small, efficient instruction-tuned baselines (1B/3B) where feasible.
 568 Comprehensive coverage of all size/variant combinations is beyond scope; we ac-
 569 knowledge this limitation and report model details and hardware requirements to
 570 inform practical deployments.

571 *7.5. Future Work*

572 To address these limitations, future work will focus on several key areas. One
 573 priority is to train the LLMs for more advanced *Declare* constructs, such as branched

574 responses. This will require generating adequate training data that accurately rep-
 575 resents these complex constructs. Additionally, developing tool support for the
 576 entire pipeline, including activity matching, will enhance the practical utility of our
 577 approach. Another important extension is to support sentences that map to multiple
 578 constraints or by managing multiple input sentences with potential inter-sentence
 579 dependencies that map to one or several output constraints. To accomplish this,
 580 post-processing steps can be helpful, e.g., splitting up activities, especially if they
 581 contain conjunctions. A more promising approach is to explicitly train the model on
 582 such scenarios. For instance, one can explore advanced techniques such as *curriculum*
 583 *learning* where the idea is to start training the model on single sentences, continue
 584 with pairs of sentences, and finally learn this task for whole paragraphs. Another idea
 585 is multi-task learning (MTL), which addresses related tasks like activity extraction,
 586 constraint extraction, sentence classification, and activity matching simultaneously.
 587 By leveraging shared representations learned from these interconnected tasks, MTL
 588 can improve model generalization and accuracy. Additionally, integrating a resource
 589 perspective can provide further insights. For instance, another task might involve
 590 mapping each activity to the actor performing it. Finally, providing target event
 591 labels as context can further improve the model's performance by offering additional
 592 information that helps identify the desired abstraction level of activities. These
 593 enhancements will make our approach more robust, versatile, and applicable to a
 594 wider range of scenarios. In addition, testing the approach on real-world datasets
 595 from various industries can provide a more comprehensive evaluation of its effec-
 596 tiveness. To address the limitation of training data, implementing automated data
 597 augmentation techniques can generate a more extensive and diverse dataset. This
 598 helps improve the model's robustness and generalization capabilities.

599 8. Conclusion

600 This work provides a systematic empirical study of LLM-based sentence-to-
 601 Declare constraint extraction, establishing foundational infrastructure for automated
 602 constraint extraction from natural language. Through comprehensive evaluation

603 across multiple model architectures and training approaches, we demonstrate that
 604 fine-tuned LLMs significantly advance the state-of-the-art, achieving 97.8% template
 605 accuracy compared to 53.8% for existing rule-based approaches. Our primary con-
 606 tributions include a benchmark dataset of 969 sentence-constraint pairs across 11
 607 Declare templates, direction-sensitive evaluation metrics that properly assess ordered
 608 constraint types, and complete reproducibility materials. The substantial performance
 609 advantages of fine-tuning over prompting approaches (99.4% vs. 56.5% template
 610 accuracy) provide clear guidance for practitioners, while competitive performance of
 611 lightweight instruction-tuned models demonstrates viable deployment across different
 612 computational constraints. Beyond technical advances, this work demonstrates clear
 613 pathways to practical decision support applications, where extracted constraints serve
 614 as machine-checkable policy abstractions enabling continuous compliance monitoring,
 615 automated conformance checking, and audit-ready governance dashboards.

616 References

- 617 [1] M. Camargo, M. Dumas, O. González, Automated discovery of business process
 618 simulation models from event logs, *Decis. Support Syst.* 134 (2020) 113284.
- 619 [2] H. van der Aa, C. Di Cicco, H. Leopold, H. A. Reijers, Extracting declarative
 620 process models from natural language, in: CAiSE, Springer, 2019, pp. 365–382.
- 621 [3] O. Ettrich, S. Stahlmann, H. Leopold, C. Barrot, Automatically identifying
 622 customer needs in user-generated content using token classification, *Decis.*
 623 *Support Syst.* 178 (2024) 114107.
- 624 [4] H. A. López, R. Strømsted, J.-M. Niyodusenga, M. Marquard, Declarative
 625 process discovery: Linking process and textual views, in: CAiSE, Springer,
 626 2021, pp. 109–117.
- 627 [5] J. C. de A. R. Gonçalves, F. M. Santoro, F. A. Baião, A case study on designing
 628 business processes based on collaborative and mining approaches, in: W. Shen,
 629 N. Gu, T. Lu, J. A. Barthès, J. Luo (Eds.), CSCWD, IEEE, 2010, pp. 611–616.

- 630 [6] F. Friedrich, J. Mendling, F. Puhlmann, Process model generation from natural
631 language text, in: CAiSE, Springer, 2011, pp. 482–496.
- 632 [7] E. V. Epure, P. Martín-Rodilla, C. Hug, R. Deneckère, C. Salinesi, Automatic
633 process model discovery from textual methodologies, in: RCIS, IEEE, 2015,
634 pp. 19–30.
- 635 [8] K. Honkisz, K. Kluza, P. Wiśniewski, A concept for generating business process
636 models from natural language description, in: KSEM, Springer, 2018, pp. 91–103.
- 637 [9] S. Sholiq, R. Sarno, E. S. Astuti, Generating BPMN diagram from textual
638 requirements, *J. King Saud Univ. Comput. Inf. Sci.* 34 (2022) 10079–10093.
- 639 [10] S. Levich, B. Lutz, D. Neumann, Utilizing the omnipresent: Incorporating
640 digital documents into predictive process monitoring using deep neural networks,
641 *Decis. Support Syst.* 175 (2023) 114043.
- 642 [11] P. Bellan, M. Dragoni, C. Ghidini, H. van der Aa, S. P. Ponzetto, Process
643 extraction from text: Benchmarking the state of the art and paving the way
644 for future challenges, 2023. [arXiv:2110.03754](https://arxiv.org/abs/2110.03754).
- 645 [12] W. V. Woensel, S. Motie, NLP4PBM: a systematic review on process extraction
646 using natural language processing with rule-based, machine and deep learning
647 methods, *Enterp. Inf. Syst.* 18 (2024).
- 648 [13] H. van der Aa, H. Leopold, M. Weidlich, Partial order resolution of event logs
649 for process conformance checking, *Decis. Support Syst.* 136 (2020) 113347.
- 650 [14] L. Ackermann, J. Neuberger, S. Jablonski, Data-driven annotation of textual
651 process descriptions based on formal meaning representations, in: CAiSE,
652 Springer, 2021, pp. 75–90.
- 653 [15] L. Quishpi, J. Carmona, L. Padró, Extracting annotations from textual
654 descriptions of processes, in: BPM, Springer, 2020, pp. 184–201.

- 655 [16] R. C. B. Ferreira, L. H. Thom, M. Fantinato, A semi-automatic approach to
 656 identify business process elements in natural language texts, in: S. Hammoudi,
 657 M. Smialek, O. Camp, J. Filipe (Eds.), ICEIS, SciTePress, 2017, pp. 250–261.
- 658 [17] C. Qian, L. Wen, A. Kumar, L. Lin, L. Lin, Z. Zong, S. Li, J. Wang, An
 659 approach for process model extraction by multi-grained text classification, in:
 660 CAiSE, Springer, 2020, pp. 268–282.
- 661 [18] J. Sànchez-Ferreres, H. van der Aa, J. Carmona, L. Padró, Aligning textual
 662 and model-based process descriptions, Data & Knowledge Engineering 118
 663 (2018) 25–40.
- 664 [19] H. van der Aa, H. Leopold, H. A. Reijers, Comparing textual descriptions to
 665 process models - the automatic detection of inconsistencies, Inf. Syst. 64 (2017)
 666 447–460.
- 667 [20] H. Van der Aa, H. Leopold, H. A. Reijers, Checking process compliance against
 668 natural language specifications using behavioral spaces, Inf. Syst. 78 (2018)
 669 83–95.
- 670 [21] M. Grohs, L. Abb, N. Elsayed, J.-R. Rehse, Large language models can accom-
 671 plish business process management tasks, in: BPM, Springer, 2023, pp. 453–465.
- 672 [22] M. Barrientos, K. Winter, J. Mangler, S. Rinderle-Ma, Verification of
 673 quantitative temporal compliance requirements in process descriptions over
 674 event logs, in: M. Indulska, I. Reinhartz-Berger, C. Cetina, O. Pastor (Eds.),
 675 CAiSE, volume 13901, Springer, 2023, pp. 417–433.
- 676 [23] H. Mustroph, M. Barrientos, K. Winter, S. Rinderle-Ma, Verifying resource
 677 compliance requirements from natural language text over event logs, in: C. D.
 678 Francescomarino, A. Burattin, C. Janiesch, S. Sadiq (Eds.), BPM, volume
 679 14159, Springer, 2023, pp. 249–265.
- 680 [24] A. Rebmann, F. D. Schmidt, G. Glavas, H. van der Aa, Evaluating the ability
 681 of llms to solve semantics-aware process mining tasks, in: ICPM, IEEE, 2024,
 682 pp. 9–16.

- 683 [25] K. Busch, T. Kampik, H. Leopold, xsemad: Explainable semantic anomaly detec-
 684 tion in event logs using sequence-to-sequence models, in: A. Marrella, M. Resinas,
 685 M. Jans, M. Rosemann (Eds.), BPM, volume 14940, Springer, 2024, pp. 309–327.
- 686 [26] J. Köpke, A. Safan, Efficient llm-based conversational process modeling, in:
 687 K. Gdowska, M. T. Gómez-López, J. Rehse (Eds.), BPM, volume 534, Springer,
 688 2024, pp. 259–270.
- 689 [27] N. Klievtsova, J. Benzin, T. Kampik, J. Mangler, S. Rinderle-Ma, Conversa-
 690 tional process modelling: State of the art, applications, and implications in
 691 practice, in: C. D. Francescomarino, A. Burattin, C. Janiesch, S. Sadiq (Eds.),
 692 BPM Forum, volume 490, Springer, 2023, pp. 319–336.
- 693 [28] M. Weidlich, J. M. E. M. van der Werf, On profiles and footprints - relational
 694 semantics for petri nets, in: S. Haddad, L. Pomello (Eds.), PETRI NETS,
 695 volume 7347, Springer, 2012, pp. 148–167.
- 696 [29] L. Barbaro, G. Varricchione, M. Montali, C. D. Ciccio, From sound
 697 workflow nets to ltlf declarative specifications by casting three spells, CoRR
 698 abs/2504.05114 (2025). [arXiv:2504.05114](https://arxiv.org/abs/2504.05114).
- 699 [30] Z. Han, C. Gao, J. Liu, J. Zhang, S. Q. Zhang, Parameter-efficient fine-tuning
 700 for large models: A comprehensive survey, CoRR abs/2403.14608 (2024).
 701 [arXiv:2403.14608](https://arxiv.org/abs/2403.14608).
- 702 [31] C. Di Ciccio, M. Montali, Declarative process specifications: reasoning,
 703 discovery, monitoring, in: Process mining handbook, Springer International
 704 Publishing Cham, 2022, pp. 108–152.
- 705 [32] M. B. Dwyer, G. S. Avruman, J. C. Corbett, Patterns in property specifications
 706 for finite-state verification, in: B. W. Boehm, D. Garlan, J. Kramer (Eds.),
 707 ICSE, ACM, 1999, pp. 411–420.
- 708 [33] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez,
 709 L. Kaiser, I. Polosukhin, Attention is all you need, in: I. Guyon, et al. (Eds.),
 710 NeurIPS, 2017, pp. 5998–6008.

- 711 [34] J. Wei, X. Wang, D. Schuurmans, M. Bosma, B. Ichter, F. Xia, E. H. Chi, Q. V.
 712 Le, D. Zhou, Chain-of-thought prompting elicits reasoning in large language
 713 models, in: S. Koyejo, S. Mohamed, A. Agarwal, D. Belgrave, K. Cho, A. Oh
 714 (Eds.), NeurIPS, 2022.
- 715 [35] T. Kojima, S. S. Gu, M. Reid, Y. Matsuo, Y. Iwasawa, Large language models
 716 are zero-shot reasoners, in: S. Koyejo, S. Mohamed, A. Agarwal, D. Belgrave,
 717 K. Cho, A. Oh (Eds.), NeurIPS, 2022.
- 718 [36] X. Wang, J. Wei, D. Schuurmans, Q. V. Le, E. H. Chi, S. Narang, A. Chowdhery,
 719 D. Zhou, Self-consistency improves chain of thought reasoning in language
 720 models, in: ICLR, OpenReview.net, 2023.
- 721 [37] E. J. Hu, Y. Shen, P. Wallis, Z. Allen-Zhu, Y. Li, S. Wang, L. Wang, W. Chen,
 722 Lora: Low-rank adaptation of large language models, in: ICLR, 2022.
- 723 [38] G. Team, T. Mesnard, C. Hardin, R. Dadashi, S. Bhupatiraju, S. Pathak,
 724 L. Sifre, M. Rivière, M. S. Kale, J. Love, et al., Gemma: Open models based
 725 on gemini research and technology, arXiv preprint arXiv:2403.08295 (2024).
- 726 [39] A. Q. Jiang, A. Sablayrolles, A. Mensch, C. Bamford, D. S. Chaplot, D. de las
 727 Casas, F. Bressand, G. Lengyel, G. Lample, L. Saulnier, L. R. Lavaud, M.-A.
 728 Lachaux, P. Stock, T. L. Scao, T. Lavril, T. Wang, T. Lacroix, W. E. Sayed,
 729 Mistral 7b, 2023. [arXiv:2310.06825](https://arxiv.org/abs/2310.06825).
- 730 [40] E. Almazrouei, H. Alobeidli, A. Alshamsi, A. Cappelli, R. Cojocaru, M. Debbah,
 731 Étienne Goffinet, D. Hesslow, J. Launay, Q. Malartic, D. Mazzotta, B. Noune,
 732 B. Pannier, G. Penedo, The falcon series of open language models, 2023.
 733 [arXiv:2311.16867](https://arxiv.org/abs/2311.16867).
- 734 [41] H. Touvron, L. Martin, K. Stone, P. Albert, A. Almahairi, Y. Babaei,
 735 N. Bashlykov, S. Batra, P. Bhargava, S. Bhosale, et al., Llama 2: Open
 736 foundation and fine-tuned chat models, arXiv preprint arXiv:2307.09288 (2023).

- 737 [42] A. Dubey, A. Jauhri, A. Pandey, A. Kadian, A. Al-Dahle, A. Letman,
 738 A. Mathur, A. Schelten, A. Yang, A. Fan, et al., The llama 3 herd of models,
 739 arXiv e-prints (2024) arXiv-2407.
- 740 [43] I. Loshchilov, F. Hutter, Decoupled weight decay regularization, in: ICLR,
 741 OpenReview.net, 2019.
- 742 [44] S. Balloccu, P. Schmidlová, M. Lango, O. Dusek, Leak, cheat, repeat: Data
 743 contamination and evaluation malpractices in closed-source llms, in: Y. Graham,
 744 M. Purver (Eds.), EACL, Association for Computational Linguistics, 2024, pp.
 745 67–93.

746 **Appendix A. Activity Matching Zero-Shot Prompting**

747 Take this list of Declare constraints containing various activities. Init(LogIn), AtMo-
 748 stOne(ResetPassword), AtLeastOne(CreateProfile), End(LogOut), Response(UserLogin,
 749 EnterMenu), Response(FileDel, ProcessFile), Precedence(UploadFile, DeleteFile), Chain-
 750 Response(ProcessingFile, ApplyChange). Match semantically equivalent activities and
 751 provide an updated list of the constraints after agreeing on one unique label per activity
 752 equivalence class.

753 **Appendix B. Activity Matching Few-Shot Prompting (excerpt)**

754 Take this list of Declare constraints containing various activities. [...] Match seman-
 755 tically equivalent activities and provide an updated list of the constraints after agreeing
 756 on one unique label per activity equivalence class. Take the following examples for
 757 which the task has been solved correctly: Input: [Response(CreateClaim, ApproveClaim),
 758 Response(ClaimCreation, SendInvoice), Init(CreatingClaim)], Output: [Response(Create-
 759 Claim, ApproveClaim), Response(CreateClaim, SendInvoice), Init(CreateClaim)] [...]

760 **Appendix C. Baseline 1: Few Shot Prompting**

761 Below is the prompt for the few-shot-prompting-based baseline used in the
 762 evaluation (cf. Subsubsection 6.5.1):

763 I am a highly intelligent AI capable of understanding and translating natural language
 764 descriptions into Declare constraints. Here are some example translations of descriptions
 765 into Declare constraints:
 766 1. Description: Breakfast must be the first to occur.
 767 Task(s): Breakfast

```

768 Declare Constraint: Init(Breakfast)
769 more shots*
770 11. Description: Cooking and Biking never occur together.
771 Task(s): Cooking, Biking
772 Declare Constraint: NotCoExistence(Cooking, Biking)
773 Given the following description, we should first detect tasks and then translate the descrip-
774 tion into a Declare constraint. The output format should exactly follow this pattern as
775 in the example descriptions. The answer should solely contain the constraint and nothing
776 else, i.e., no additional description. Now, let's translate the following description into a
777 Declare constraint:
778 12. Description: *description*
779

```

780 Appendix D. Baseline 2: Few Shot Chain-of-Thought (CoT) Prompting

781 Below is the excerpt of the prompt for the few-shot-CoT-prompting-based base-
 782 line used in the evaluation (cf. Subsubsection 6.5.1). Note that we employ a specific
 783 syntax for this prompt. The instructional version of Google's Gemma 7b-bit model
 784 we chose for inference utilizes control tokens to enhance the clarity and structure
 785 of conversations. This syntax includes tokens for indicating the user's turn (`user`),
 786 the model's turn (`model`), the start of a dialogue turn (`<start_of_turn>`), and the
 787 end of a dialogue turn (`<end_of_turn>`). This clear differentiation of roles and turns
 788 improves the model's performance by providing explicit context and structure.

```

789 <start_of_turn>user
790 You are a highly intelligent AI capable of understanding and translating natural language de-
791 scriptions into Declare constraints. Here are definitions of some common Declare templates:
792 1. Definition: "Task A must be the first to occur."
793 Declare: Init(A)
794 more definitions*
795 11. Definition: "Task A and Task B never occur together."
796 Declare: NotCoExistence(A, B)
797 Think and write your step-by-step reasoning before responding.
798 Now, translate the following description into a Declare constraint: "They have to create
799 the claim, before they can approve it."
800 <end_of_turn>
801 <start_of_turn>model
802 Let's think step by step. In this example, we find two tasks, namely, creating a claim and
803 approving it. We, therefore, define CreateClaim and ApproveClaim as tasks. We define
804 Task A = CreateClaim and Task B = ApproveClaim because of the temporal sequence
805 of these tasks. According to the description, CreateClaim (Task A) is a condition for

```

806 ApproveClaim (Task B). This matches the definition of a Precedence constraint (If Task
 807 B occurs, then Task A must occur before).
 808 Thus, the resulting constraint is: Precedence(CreateClaim, ApproveClaim)
 809 <end_of_turn>
 810 more shots with chain-of-thoughts*
 811 <start_of_turn>user
 812 Now, translate the following description into a Declare constraint: "Booking the venue
 813 should only happen once."
 814 <end_of_turn>
 815 <start_of_turn>model
 816 Let's think step by step. In this example, we find one task, namely booking the venue. We,
 817 therefore, define BookVenue as the task. We define Task A = BookVenue. According to
 818 the description, BookVenue (Task A) should only happen once. This matches the definition
 819 of an AtMostOne constraint (Task A occurs at most once).
 820 Thus, the resulting constraint is: AtMostOne(BookVenue)
 821 <end_of_turn>
 822 <start_of_turn>user
 823 Now, let's translate the following description into a Declare constraint: description.
 824 <end_of_turn>
 825 <start_of_turn>model
 826 Let's think step by step.

Extracting Declarative Constraints for Process Modeling From Natural Language Descriptions with Large Language Models

Gyunam Park^{a,b}, Julian Kofferath^c, Minsu Cho^{d,*}

^a*Eindhoven University of Technology, Eindhoven, 5612 AZ, the Netherlands*

^b*Fraunhofer Institute for Applied Information Technology FIT, Sankt Augustin, 53757, Germany*

^c*RWTH Aachen University, Aachen, 52074, Germany*

^d*School of Information Convergence, Kwangwoon University, Seoul, 01897, South Korea*

Acknowledgements.

This work was supported by the Ministry of Education of the Republic of Korea and the National Research Foundation of Korea (NRF-2024S1A5A8024081) and the Research Grant of Kwangwoon University in 2024.

*Corresponding author

Email addresses: g.park@tue.nl (Gyunam Park), julian.kofferath@rwth-aachen.de (Julian Kofferath), mcho@kw.ac.kr (Minsu Cho)

Highlights

1. LLM-based approach for automated extraction of declarative process models.
2. Novel dataset with 969 labeled text inputs covering 11 Declare templates.
3. Enhanced generalization and accuracy over existing rule-based NLP techniques.

Biographical Note

Gyunam Park is an Assistant Professor in computer science at Eindhoven University of Technology. Prior to this, he stayed at Fraunhofer FIT as a research group leader. He completed his Ph.D. in Computer Science at RWTH Aachen University. His research interests include process mining, machine learning, and the application of large language models. Dr. Park has published extensively in high-impact journals, including *Decision Support Systems*, *Expert Systems with Applications (ESWA)*, *Computers in Industry*, and *Engineering Applications of Artificial Intelligence*. His work contributes to advancing the understanding of data-driven decision support and intelligent process automation in various industrial contexts.

Julian Kofferath received his B.Sc. degree in Computer Science from RWTH Aachen University, where he is currently pursuing his M.Sc. in Computer Science. His research interests encompass data science, process mining, and the application of large language models. Julian has been involved in a research project funded by the Federal Ministry of Education and Research in Germany.

Minsu Cho received his Ph.D. degree in management engineering at UNIST (Ulsan National Institute of Science and Technology) in 2018. He is now an associate professor with the School of Information Convergence at Kwangwoon University. His research interests include process mining, natural language processing, business analytics, and industrial artificial intelligence. He has published more than 30 scientific papers in several top-level venues such as *Decision Support Systems*, *Expert Systems with Applications*, *Journal of Biomedical Informatics*, *International Journal of Medical Informatics*, etc.

Declaration of interests

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

The authors declare the following financial interests/personal relationships which may be considered as potential competing interests: